

Relational Programming in miniKanren

Relational Programming in miniKanren

William E. Byrd

© 2014 William E. Byrd

Typeset by the author in X_YL^AT_EX, using Tufte-Style Book (<http://www.LaTeXTemplates.com>).

June 18, 2014 version.

The latest version of this book, along with all files necessary to typeset it, can be found at:

<https://github.com/webyrd/relational-programming-in-miniKanren>



This work is licensed under a Creative Commons Attribution 4.0 International License. (CC BY 4.0)

<http://creativecommons.org/licenses/by/4.0/>

For my H211 students:

Indiana University, Fall 2010 & 2011, and Team pw0ni3.

Learning with always trumps learning from.

—Woodie Flowers

Contents

Preface ix

Introduction 1

Relational Interpreters 3

Conclusion 5

Bibliography 7

Preface

This book is about *relational programming*. Just as functional programs model mathematical functions, relational programs model mathematical relations. This relational view of programming erases the distinction between a program’s inputs and outputs. The resulting code is more flexible than the functional, imperative, or object-oriented equivalents, often in delightfully surprising and enlightening ways. For example, we can run a relational program *backwards*, determining which “input” values correspond to a given “output” value.

Relational programming can be seen as an especially pure form of traditional (Prolog-style) logic programming, in which we commit completely to the ideal of writing programs as relations. We are willing to trade efficiency and expressiveness in favor of increased flexibility for the programs we *can* express. In order to express more interesting programs, and to regain efficiency, we will augment our core relational language with an ever-larger set of *constraint operators*. For this reason, our approach is closer to that of constraint logic programming than to traditional logic programming. In contrast with other forms of logic programming, we will always emphasize the *relational* aspect of our programs. Our motto for this book shall be:

A program that doesn’t run backwards isn’t worth writing.

The two main challenges we will face when writing programs as relations are *avoiding divergence* (looping forever) and *expressing negation*, which turn out to intimately connected.

Audience

This book is written for intermediate-to-advanced programmers, computer science students, and researchers. For this book, *intermediate* means that you are comfortable writing simple recursive procedures in a functional programming language, such as Scheme, Racket, Clojure, Lisp, ML, or Haskell. I also assume you have a reading knowledge of Scheme. No knowledge of relational programming, logic programming, or programming language theory is required.

This means avoiding all “extra-logical” operators, such as Prolog’s `cut (!)`, `assert` and `retract`, `is`, and `copy_term`.

Pure Prolog is actually a special case of constraint logic programming, over the domain of trees (*CLP(tree)*).

With apologies to Alan Perlis.

If Scheme is the best language for learning recursion, and Haskell for learning about monads, then miniKanren is the best language for understanding the Halting Problem.

I have attempted to deliver [these lectures] in a spirit that should be recommended to all students embarking on the writing of their PhD theses: imagine that you are explaining your ideas to your former smart, but ignorant, self, at the beginning of your studies!

—Richard P. Feynman
The Feynman Lectures on Computation

If you want to learn about relational programming, but are new to programming, Dan Friedman, Oleg Kiselyov, and I have written a book just for you, called *The Reasoned Schemer*¹. In that book we assume you are familiar with the material in *The Little Schemer*², which is a very gentle introduction to recursion and functional programming.

You might also benefit from *The Little Schemer* if you are an experienced programmer but feel uncomfortable with recursion. More traditional texts on functional programming in Scheme include *Scheme and the Art of Programming*³, the first edition of *How to Design Programs*⁴, and the classic *Structure and Interpretation of Computer Programs*⁵.

If you are an experienced functional programmer, but do not know Scheme, the beginning of *Structure and Interpretation of Computer Programs* should get you up to speed, while *The Scheme Programming Language, Fourth Edition*⁶ describes the language in detail.

The Language

The relations in this book are written in *miniKanren*, a language designed specifically for relational programming. Other languages that appear well-suited to relational programming include the logic programming languages λ Prolog⁷ and Gödel⁸, and the Twelf specification and proof language.

miniKanren was first implemented in Scheme; over the past eight years miniKanren has been implemented in Racket, Clojure, Ruby, Python, and many other “host” languages. There are dozens of experimental miniKanren implementations in Scheme alone, and several other host languages have more than one miniKanren implementation. Like Lisp and Scheme, miniKanren is really a *family* of related languages.

This diversity of implementations and host languages has its advantages. Interest in miniKanren has largely been driven by the popularity of David Nolen’s excellent `core.logic` Clojure library. Claire Alvis continues to add advanced constraint solving features to the `cKanren` (“constraint Kanren”) Racket library. Other researchers and hackers are experimenting with their own language extensions and optimizations, greatly accelerating the evolution of miniKanren. Unfortunately, this diversity poses a dilemma for both the author and readers of this book: which host language and implementation to use?

Since Scheme is the language I know best, and since most academic work on miniKanren has used Scheme as the host language, I have decided to write this book using Scheme. Anyone interested in “porting” the book to Racket, Clojure, or another host language may do so with my blessing.

¹ D. P. Friedman, W. E. Byrd, and O. Kiselyov. *The Reasoned Schemer*. MIT Press, Cambridge, MA, 2005

² D. P. Friedman and M. Felleisen. *The Little Schemer* (4th ed.). MIT Press, Cambridge, MA, 1996

³ G. Springer and D. P. Friedman. *Scheme and the Art of Programming*. MIT Press, Cambridge, MA, 1989

⁴ M. Felleisen, R. B. Findler, M. Flatt, and S. Krishnamurthi. *How to Design Programs: An Introduction to Programming and Computing*. MIT Press, Cambridge, MA, first edition, 2001

(full text at <http://htdp.org/>)

⁵ H. Abelson and G. J. Sussman. *Structure and Interpretation of Computer Programs*. MIT Press, Cambridge, MA, 2nd edition, 1996

(full text at <http://mitpress.mit.edu/sicp/full-text/book/book.html>)

⁶ R. K. Dybvig. *The Scheme Programming Language, 4th Edition*. The MIT Press, 4th edition, 2009

(full text at <http://www.scheme.com/tspl4/>)

⁷ D. Miller and G. Nadathur. *Programming with Higher-Order Logic*. Cambridge University Press, New York, NY, 2012

⁸ P. M. Hill and J. W. Lloyd. *The Gödel programming language*. MIT Press, 1994

Twelf website: <http://twelf.org/>
Links to these implementations, and many other miniKanren resources, can be found at the official miniKanren website: <http://minikanren.org/>

`core.logic`:

<https://github.com/clojure/core.logic>

`cKanren` in Racket:

<https://github.com/calvis/cKanren>

Goal

The goal of this book is to show how to effectively use miniKanren to write relational programs. To this end, the book presents a variety of non-trivial miniKanren relations, written in what I consider to be idiomatic style, and shows how these relations were derived. Many of these examples draw from academic papers on miniKanren, and focus on programming language theory (interpreters, type inferencers, etc.). Other examples, such as finite state machines, should be immediately understandable by a wider audience of programmers. To make the book as accessible as possible, all of these concepts are explained either in the main text or in appendices. This book can therefore be used as an informal introduction to programming language concepts.

Running the Code

All of the source code in the book can be found at:

<https://github.com/webyrd/relational-programming-in-miniKanren/tree/master/src>

All code has been tested under Vicare Scheme, and should run with little or no modification under other implementations that are compatible with the R^5RS or R^6RS Scheme standards.

Margin Notes

This book is typeset in the style of Edward Tufte’s magnificent and beautiful *The Visual Display of Quantitative Information*⁹. I share Tufte’s love of margin notes, and use them in this book to help solve the problem of addressing readers with widely varying knowledge of computer science and programming. To make the book as accessible as possible, in the main text I assume the reader is the hypothetical intermediate-level programmer or student described in the *Audience* section above. In the margin notes, however, anything goes, and I reserve the right to geek out whenever necessary (or unnecessary).

Exercises

Typographic Conventions

Acknowledgements

William E. Byrd
Salt Lake City, Utah
June 2014

In this sense the book emulates Benjamin Pierce’s outstanding *Types and Programming Languages* (TAPL), which is about types, but also contains excellent introductions to other programming languages related topics:

B. C. Pierce. *Types and Programming Languages*. MIT Press, Cambridge, MA, 2002

Vicare Scheme, developed by Marco Maggi, is a fork of my friend Abdulaziz Ghuloum’s Ikarus Scheme: <http://marcomaggi.github.io/vicare.html>
 R^5RS Standard: <http://www.schemers.org/Documents/Standards/R5RS/>
 R^6RS Standard: <http://www.r6rs.org/>

This book is set using the “Tufte-Style Book” \LaTeX style, freely available from <http://www.LaTeXTemplates.com>

⁹ E. R. Tufte. *The Visual Display of Quantitative Information*. Graphics Press, Cheshire, CT, 1986

Another great lover of both margin notes and footnotes was David Foster Wallace (1962–2008). DFW’s mastery of notes is best demonstrated by his essay, “Host,” in:

D. F. Wallace. *Consider the Lobster and Other Essays*. Little, Brown and Co., 2005

Wallace also loved *marginalia*; the Harry Ransom Center’s DFW collection includes heavily annotated books from Wallace’s personal library: <http://www.hrc.utexas.edu/press/releases/2010/dfw/books/>.

Introduction

gl hf!

—Greg “IdrA” Fields

Relational programs generalize functional programs, in that they do not distinguish between the “input” arguments passed to a function and the “output” result returned by that function. For example, consider a two-argument variant of Scheme’s addition function, restricted to natural numbers: $(+ \ 3 \ 4) \Rightarrow 7$. A relational version of addition, $+^o$, takes three arguments: $(+^o \ 3 \ 4 \ z)$, where z is a *logic variable* representing the result of adding the first two arguments of $+^o$. In this case z is associated with 7. More interestingly, we can write $(+^o \ 3 \ y \ 7)$, which associates y with 4; our addition relation also performs subtraction. We can also write $(+^o \ x \ y \ 7)$, which associates x and y with all pairs of natural numbers that sum to 7; $+^o$ produces multiple answers, including $x = 3$ and $y = 4$, and $x = 0$ and $y = 7$. Finally, we can write $(+^o \ x \ y \ z)$, which enumerates all triples of natural numbers (x, y, z) such that $x + y = z$; $+^o$ produces infinitely many answers. Informally, we say that the call $(+^o \ 3 \ 4 \ z)$ runs the $+^o$ relation “forwards,” while the calls $(+^o \ 3 \ y \ 7)$, $(+^o \ x \ y \ 7)$, and $(+^o \ x \ y \ z)$ run “backwards.”

We will see that the remarkable flexibility of the $+^o$ relation is exhibited by more complex relations, such as interpreters, type inferencers, and finite-state machines. For example, we will write a relational interpreter for a subset of Scheme, $eval^o$. Running $eval^o$ forwards, the call $(eval^o \ '((\lambda (x) (+ \ x \ 2)) \ 4) \ val)$ associates val with 6. Running backwards is more interesting: $(eval^o \ exp \ '6)$ generates legal Scheme expressions that *evaluate* to 6, while $(eval^o \ exp \ exp)$ generates *quines*, which are Scheme expressions that evaluate to themselves.

The *natural numbers* are the non-negative integers: 0, 1, ...

Here we are taking a notational liberty, as $+^o$ expects 3 and 4 to be represented as *binary, little-endian* lists: (1 1) and (0 0 1), respectively. Zero is uniquely represented as the empty list, (). To ensure a unique representation of each number, lists may not end with the digit 0. This numeric representation is extremely flexible, since the lists can contain logic variables—for example, the list '(1 . , x) represents any odd natural number, while '(0 . , x) represents any positive even natural.

We can also perform relational arithmetic on built-in Scheme numbers, using *Constraint Logic Programming over Finite Domains*, or CLP(FD); as we will see, CLP(FD) is faster, but less general, than $+^o$ and friends ($*^o$, $/^o$, etc.).

As might be expected, $eval^o$ will interpret the quoted list $(+ \ x \ 2)$ as the call $(+^o \ x \ 2 \ val)$, where the logic variable x is associated with 4.

Relational Interpreters

We will never run out of things to program as long as there is a single program around.

—Alan J. Perlis
Epigrams on Programming, #100

Conclusion

G.G.

—Sean “Day[9]” Plott

Bibliography

H. Abelson and G. J. Sussman. *Structure and Interpretation of Computer Programs*. MIT Press, Cambridge, MA, 2nd edition, 1996.

R. K. Dybvig. *The Scheme Programming Language, 4th Edition*. The MIT Press, 4th edition, 2009.

M. Felleisen, R. B. Findler, M. Flatt, and S. Krishnamurthi. *How to Design Programs: An Introduction to Programming and Computing*. MIT Press, Cambridge, MA, first edition, 2001.

D. P. Friedman and M. Felleisen. *The Little Schemer (4th ed.)*. MIT Press, Cambridge, MA, 1996.

D. P. Friedman, W. E. Byrd, and O. Kiselyov. *The Reasoned Schemer*. MIT Press, Cambridge, MA, 2005.

P. M. Hill and J. W. Lloyd. *The Gödel programming language*. MIT Press, 1994.

D. Miller and G. Nadathur. *Programming with Higher-Order Logic*. Cambridge University Press, New York, NY, 2012.

A. J. Perlis. Special feature: Epigrams on programming. *SIGPLAN Not.*, 17(9):7–13, Sept. 1982. URL <http://doi.acm.org/10.1145/947955.1083808>.

B. C. Pierce. *Types and Programming Languages*. MIT Press, Cambridge, MA, 2002.

G. Springer and D. P. Friedman. *Scheme and the Art of Programming*. MIT Press, Cambridge, MA, 1989.

E. R. Tufte. *The Visual Display of Quantitative Information*. Graphics Press, Cheshire, CT, 1986.

D. F. Wallace. *Consider the Lobster and Other Essays*. Little, Brown and Co., 2005.